

STAN49 - Assignment 2

William Nordansjö

March 13, 2025

Introduction

This assignment will consist of two parts covering two different aspects of textual data analysis. The first part relates to word embeddings, a skip-gram method will be employed in order to create word embeddings, these will then be compared to a pre-trained set of embeddings in training a deep neural network to perform classification. The second part of this assignment is more free form, and the task is to identify the odd one out of a collection of French plays. This will be done using an unsupervised learning approach by application of clustering and dimension reduction methods.

Task 1: AG's News Topic Classification

In this task I aim to classify the topic of 7600 news articles by training a deep learning neural network on 120 000 other articles. The documents are courtesy of Gulli (2015) and consists of news articles from the four categories: *World*, *Sports*, *Business*, and *Sci/Tech*. To help me in this endeavor I will use two different word embeddings: A locally trained model using skip-grams with *Word2Vec*, in addition the embeddings of pre-trained model called GloVe will be used as a benchmark when training and fitting the neural network. Word embeddings assign a value to a word in a vocabulary based on its similarity to neighboring words, and can as such be a very powerful tool in natural language processing by condensing a large amount of information, including context and a proxy for sentiment of a word in a single value.

The first word embeddings is thus the skip-gram embeddings. In order to obtain these, the dataset was loaded into python, the articles content was concatenated with their titles in order to form a dataframe of documents with a category and some text. This text was preprocessed in the usual fashion; predefined

english stopwords was removed, the sentences was forced into lowercase, dots and commas was removed and finally each word was converted into a token. After this, each word in the vocabulary was assigned an index, this was then flipped so that each indexed value was paired with a word. The result of the preprocessing can be found in table 1.

Table 1: Example Text and Sequence after Preprocessing.

Text	Sequence of indices
wall st bears claw back black reuters reuters ...	[31825, 26071, 69558, 76225, 77386, 68455, 287...]
carlyle looks toward commercial aerospace reut...	[23682, 21340, 65845, 56306, 80426, 28778, 287...]

The next step in acquiring the embeddings is to generate the training data required to fit the *Word2Vec* function. The training data consist of positive skip-gram pairs which are true empirical pairs of words which creates a target and negative skip-grams which are out of sample pairs or false context. The *Word2Vec* function is then applied to the training data and thus creates the word embedding vector by trying to predict the context of a given word. Applied on the data created above, the function achieved a very high degree of accuracy as can be seen in table 2. After training we now have our locally trained, skip-gram, 100 dimensional word embeddings.

Table 2: Model Evaluation Results

Metric	Value
Accuracy	0.9962
Loss	0.0221

The other set of word embeddings is as mentioned previously transferred from a pre-trained GloVe model. A GloVe model is in contrast to the skip-gram model trained on global co-occurrences between word pairs, this means that it should be able to capture larger trends in the vocabulary. The pre-trained word vector was extracted from Pennington et. al (2014) and in order to align with the skip-gram model, the 100 dimension version was used. The extracted data file contained 400000 word vectors which was matched with words in the vocabulary of our news article dataset.

The model architecture chosen is based on the Bidirectional Gated Recurrent Unit, (BiGru) and consist of an embedding layer, three recurrent units and one dense layer before a final dense softmax output layer. Each BiGru and dense layer has an accompanying dropout regularization layer. This architecture will be applied two times the news data in a multiclass setting in order to predict which news category a given article is about, both times the model will be tuned over overall dropout rate but also unit size of the recurrent and dense layer. The main difference of the two applications will be which word embeddings that are used.

The result of the tuning of the skip-gram based model can be found in table 3. And the performance of the model can be seen in table 5 as well as in figure 1. Similarly, the tuning made using the GloVe embeddings is stored in table 4, overall classification performance metrics can be found in table 6 and finally the classifications themselves are visualized in figure 2.

Table 3: Optimal Hyperparameters for Skip-gram Model Found by Keras Tuner

Hyperparameter	Optimal Value
GRU Unit 1	128
GRU Unit 2	256
GRU Unit 3	128
Dense Units	64
Dropout Rate	0.2
Batch Size	512
Optimizer	Adam

Table 4: Optimal Hyperparameters for GloVe Model Found by Keras Tuner

Hyperparameter	Optimal Value
GRU Unit 1	192
GRU Unit 2	128
GRU Unit 3	128
Dense Units	128
Dropout Rate	0.2
Batch Size	512
Optimizer	Adam

Overall, both models performed well by achieving a high degree of accuracy, with the Skip-Gram model correctly classifying 84% of the sequences and the GloVe model scored even higher at 92%. Seeing however that this is a multi-class classification scenario, simply looking at accuracies only gets you so far. In the classification reports, we can get a sense of how successful the models were in classifying specific categories. For both models, *Sports* was by far the easiest category to classify while *Business* was the hardest. Even more interesting, by looking at the confusion matrices, we can see how the categories interact with each other. Here we can see that both models had trouble discerning between *Business* and *Sci/Tech*, probably because many of the finance related news articles are related to scientific advances. Given that we have designed our experiment correctly, the only difference between the models are the word embeddings. Thus we can accredit the difference in performance the GloVe models capability of capturing global trends, and while the Skip-Gram generated embeddings produced a competent model, it is restricted by its local nature.




Figure 1: Confusion Matrix of Classifications Made by the Skip-gram Model




Figure 2: Confusion Matrix of Classifications Made by the GloVe Model

Table 5: Classification Report: Precision, Recall, and F1-Score

Class	Precision	Recall	F1-Score	Support
World	0.88	0.84	0.86	5956
Sports	0.94	0.90	0.92	6058
Business	0.79	0.87	0.83	5911
Sci/Tech	0.82	0.83	0.83	6075
Accuracy	0.86 (24000 samples)			
Macro Avg	0.86	0.86	0.86	24000
Weighted Avg	0.86	0.86	0.86	24000

Table 6: Classification Report: Precision, Recall, and F1-Score

Class	Precision	Recall	F1-Score	Support
World	0.94	0.91	0.93	5956
Sports	0.96	0.99	0.97	6058
Business	0.89	0.89	0.89	5911
Sci/Tech	0.90	0.90	0.90	6075
Accuracy	0.92 (24000 samples)			
Macro Avg	0.92	0.92	0.92	24000
Weighted Avg	0.92	0.92	0.92	24000

Task 2: Identify the Odd Tragedy

In this section, the task is to identify the odd one out from a collection of French plays, tragedies to be specific. This collection makes up a corpus of the speaking lines of ten plays written in French. Nine of the plays are written by one author and one play is written by a different author. In aid of our mission to root which of the plays it is, we have the statistical field of unsupervised learning, of which we will use three methods: K-means clustering, Hierarchical clustering and finally PCA or Principal Component Analysis.

First of however, we have to pre-process our text into a more usable format. This will be done by first removing dots and commas and coercing lowercase lettering, before removing *French* stopwords (relying on the nltk package to make those distinctions) and turning the words into tokens. Finally the documents were converted to tf-idf vectors, meaning that each term was weighted in relation to how rare it is in its document.

The first clustering approach applied to the text was k-means clustering. This is an iterative EM-, or Expectation Maximization method which tries to sort the data into a given amount of groups, by calculating the mean or centroid of those groups and stepwise tries to find observations that are closest to this centroid. In this case we are interested in finding two groups, one of plays written by the main




Figure 3: Hierarchical Clustering Dendrogram

author and the other consisting of the play written by the other author. The package used courtesy of sklearn utilizes the k-means++ initialization algorithm which places the initial clusters evenly among the data and then iteratively shifts the centroid until two distinct groups are formed. The result can be found in table 7, and shows that the algorithm found two groups, one with 8 plays and only containing plays 6 and 9.

Table 7: Cluster Assignments of Plays

Play	Cluster	Play	Cluster
Play 01	1	Play 06	0
Play 02	1	Play 07	1
Play 03	1	Play 08	1
Play 04	1	Play 09	0
Play 05	1	Play 10	1

The second clustering approach is hierarchical clustering which attempts to group the observations based upon some measure of distance. There are many of these measures, often called linkages, to choose from but in this case *Ward's* method is employed. *Ward* linkage tries to find clusters such that they minimize within cluster variance, starting from one cluster and then splitting and grouping observations until each is its own cluster. This process is visualized in the dendrogram in figure 3, and interestingly, when cutting the diagram horizontally such that two clusters are formed, these are identical to the clusters formed by the k-means algorithm. If however the dendrogram was cut into three clusters, play number 7 would make a cluster of its own.

The final method employed in this analysis is Principal Component Analysis (PCA). PCA is a widely used dimensionality reduction technique that transforms high-dimensional data into a lower-dimensional space




Figure 4: PCA Projection of Plays

while preserving as much variance as possible. The method works by computing the eigenvectors and eigenvalues of the data's covariance matrix, where each principal component corresponds to an eigenvector that maximizes the variance along that direction. These principal components are orthogonal to each other and ordered by their ability to explain variance in the dataset. In figure 4 the first two principal components are plotted against each other and are assigned to one of two groups. The pattern seen previously is again repeated here where plays 6 and 9 are separate when looking at two groups but when a third is considered play number 7 again stands out.

From our analysis we have found compelling evidence for the play written by the other author. First of all, when only looking at two clusters, the suspect is hidden among its stylized neighbors and its true nature is only revealed by examining a three cluster scenario. The working theory is thus that we have two authors, but also two types or styles of plays. Our previous knowledge is that all of the plays are tragedies, but apparently these two types of tragedies are distinct to our models. These styles are nigh on impossible to discern without knowledge of either French literary style or even French as a language, but perhaps there is a different tone to it. This leads us to conclude: There are two styles, the first contains two plays written by our main author, and the second contains 8 plays, 7 of which are written by the same, and one written by the other writer. As such, play number 7 is the odd tragedy.

Conclusion

In this assignment, two distinct textual analysis tasks were completed using deep-learning and unsupervised learning techniques. The first task was classifying news topics, where a deep-learning model was trained on locally generated skip-gram embeddings and compared to a similar model using pre-trained GloVe embeddings instead. The latter outperformed the former, highlighting the benefits of capturing global word relationships. In the second task an odd French tragedy was supposed to be found. Here k-means, hierarchical clustering, and PCA consistently found that plays 6 and 9 differed from the rest, but examining a three-cluster scenario revealed play 7 to be a separate cluster. This suggests the existence of two stylistic groups, with play 7 standing apart as the odd one out.

References

Gulli, A. 2015. *AG's corpus of news articles*, Link: http://groups.di.unipi.it/~gulli/AG_corpus_of_news_articles.html

Jeffrey Pennington, Richard Socher, and Christopher D. Manning. 2014. *GloVe: Global Vectors for Word Representation*, Link: <https://nlp.stanford.edu/projects/glove/>

Appendix A: Python Code

```
1
2 import numpy as np
3 import pandas as pd
4 import os
5 import random
6 import re
7 import tensorflow as tf
8 from tensorflow.keras.preprocessing.sequence import pad_sequences
9 from sklearn.model_selection import train_test_split
10 from tensorflow.keras.models import Sequential
11 from tensorflow.keras.layers import Embedding, SimpleRNN, LSTM, GRU, Dense,
    Dropout, Input, Bidirectional
12 from tensorflow.keras.callbacks import EarlyStopping
13 import keras_tuner as kt
14 import matplotlib.pyplot as plt
15 import seaborn as sns
16 from sklearn.metrics import confusion_matrix
17 import tqdm
18 from tensorflow.keras import layers
19 from collections import Counter
20 from nltk.tokenize import word_tokenize
21 from nltk.corpus import stopwords
22 from nltk.stem import WordNetLemmatizer
23 import nltk
24
25 #Data
26 column_names = ["Category", "Title", "Description"]
```

```

27
28 train = pd.read_csv("/content/AG_train.csv", names = column_names, header = None
    )
29 test = pd.read_csv("/content/AG_test.csv", names = column_names, header = None)
30
31 train["Text"] = train["Title"] + " " + train["Description"]
32 test["Text"] = test["Title"] + " " + test["Description"]
33 train.drop(columns=["Title", "Description"], inplace=True)
34 test.drop(columns=["Title", "Description"], inplace=True)
35
36 #Preprocessing
37 nltk.download("stopwords")
38 stop_words = set(stopwords.words("english"))
39
40 def preprocess_text(sentence):
41     sentence = re.sub(r"[^\w\s]", "", sentence)
42     tokens = sentence.lower().split()
43     tokens = [word for word in tokens if word not in stop_words]
44     return " ".join(tokens)
45
46 train["Text"] = train["Text"].apply(preprocess_text)
47 test["Text"] = test["Text"].apply(preprocess_text)
48
49 #Creating vocabulary
50 all_tokens = []
51 for text in train["Text"]:
52     all_tokens.extend(text.split()) # Tokenizing by space
53
54 # Get unique words
55 unique_tokens = set(all_tokens)
56 print(f"Total unique words: {len(unique_tokens)}")
57
58 # Initialize vocab with a padding token
59 vocab = {"<pad>": 0} # Start indexing from 1
60 index = 1
61
62 # Assign an index to each unique word
63 for token in unique_tokens:

```

```

64     vocab[token] = index
65     index += 1
66
67 vocab_size = len(vocab) # Number of unique words including <pad>
68 print(f"Vocabulary size: {vocab_size}")
69
70 # Reverse dictionary to get index word mapping
71 inverse_vocab = {index: token for token, index in vocab.items()}
72
73 # Convert dataset text into sequences of word indices
74 train["Sequences"] = train["Text"].apply(lambda x: [vocab[word] for word in x.
75     split() if word in vocab])
76
77 test["Sequences"] = test["Text"].apply(lambda x: [vocab[word] for word in x.
78     split() if word in vocab])
79
80 # Show an example
81 print(train[["Text", "Sequences"]].head())
82
83 # Generates skip-gram pairs with negative sampling for a list of sequences
84 # (int-encoded sentences) based on window size, number of negative samples
85 # and vocabulary size.
86
87 def generate_training_data(sequences, window_size, num_ns, vocab_size, seed):
88     # Elements of each training example are appended to these lists.
89     targets, contexts, labels = [], [], []
90
91     # Build the sampling table for 'vocab_size' tokens.
92     sampling_table = tf.keras.preprocessing.sequence.make_sampling_table(
93         vocab_size)
94
95     # Iterate over all sequences (sentences) in the dataset.
96     for sequence in tqdm.tqdm(sequences):
97
98         # Generate positive skip-gram pairs for a sequence (sentence).
99         positive_skip_grams, _ = tf.keras.preprocessing.sequence.skipgrams(
100             sequence,
101             vocabulary_size=vocab_size,
102             sampling_table=sampling_table,
103             window_size=window_size,

```

```

99         negative_samples=0,
100         seed=seed)
101
102     # Iterate over each positive skip-gram pair to produce training examples
103     # with a positive context word and negative samples.
104     for target_word, context_word in positive_skip_grams:
105         context_class = tf.expand_dims(
106             tf.constant([context_word], dtype="int64"), 1)
107         negative_sampling_candidates, _, _ = tf.random.
108             log_uniform_candidate_sampler(
109                 true_classes=context_class,
110                 num_true=1,
111                 num_sampled=num_ns,
112                 unique=True,
113                 range_max=vocab_size,
114                 seed=seed,
115                 name="negative_sampling")
116
117     # Build context and label vectors (for one target word)
118     context = tf.concat([tf.squeeze(context_class,1),
119         negative_sampling_candidates], 0)
120     label = tf.constant([1] + [0]*num_ns, dtype="int64")
121
122     # Append each element from the training example to global lists.
123     targets.append(target_word)
124     contexts.append(context)
125     labels.append(label)
126
127     return targets, contexts, labels
128
129 #targets, contexts, labels = generate_training_data(
130 #    sequences=train["Sequences"],
131 #    window_size=3,
132 #    num_ns=4,
133 #    vocab_size=vocab_size,
134 #    seed=407)
135
136 #targets = np.array(targets)

```

```

135 #contexts = np.array(contexts)
136 #labels = np.array(labels)
137
138 #print('\n')
139 #print(f"targets.shape: {targets.shape}")
140 #print(f"contexts.shape: {contexts.shape}")
141 #print(f"labels.shape: {labels.shape}")
142
143 BATCH_SIZE = 200
144 BUFFER_SIZE = 5000
145 AUTOTUNE = tf.data.AUTOTUNEa
146
147 dataset = tf.data.Dataset.from_tensor_slices(((targets, contexts), labels))
148 dataset = dataset.shuffle(BUFFER_SIZE).batch(BATCH_SIZE, drop_remainder=True)
149 dataset = dataset.cache().prefetch(buffer_size=AUTOTUNE)
150
151 #Word2Vec Model
152 class Word2Vec(tf.keras.Model):
153     def __init__(self, vocab_size, embedding_dim):
154         super(Word2Vec, self).__init__()
155         self.target_embedding = layers.Embedding(vocab_size,
156                                                    embedding_dim,
157                                                    input_length=1,
158                                                    name="w2v_embedding")
159         self.context_embedding = layers.Embedding(vocab_size,
160                                                    embedding_dim,
161                                                    input_length=num_ns+1)
162
163     def call(self, pair):
164         target, context = pair
165         # target: (batch, )
166         # context: (batch, context)
167         if len(target.shape) == 2:
168             target = tf.squeeze(target, axis=1)
169         # target: (batch,)
170         word_emb = self.target_embedding(target)
171         # word_emb: (batch, embed)
172         context_emb = self.context_embedding(context)

```

```

173     # context_emb: (batch, context, embed)
174     dots = tf.einsum('be,bce->bc', word_emb, context_emb) # 'be, bce ->'
           indicates the output shape
175     # dots: (batch, context)
176     return tf.nn.softmax(dots)
177
178 num_ns=4
179 embedding_dim = 100
180 word2vec = Word2Vec(vocab_size, embedding_dim)
181 word2vec.compile(optimizer='adam',
182                 loss="CategoricalCrossentropy",
183                 metrics=['accuracy'])
184
185 #weights = word2vec.get_layer('w2v_embedding').get_weights()[0]
186
187 MAX_SEQUENCE_LENGTH = 100
188 NUM_CLASSES = len(train["Category"].unique())
189
190 train_labels = train["Category"].values -1
191
192 padded_sequences = pad_sequences(train["Sequences"], maxlen=MAX_SEQUENCE_LENGTH,
           padding="post")
193
194 X_train, X_test, y_train, y_test = train_test_split(padded_sequences,
           train_labels, test_size=0.2, random_state=42)
195 y_train = tf.keras.utils.to_categorical(y_train, num_classes=NUM_CLASSES)
196 y_test = tf.keras.utils.to_categorical(y_test, num_classes=NUM_CLASSES)
197
198 #Model 1: Deep learning model using skip-gram word embedding
199
200 def model_skip(hp):
201     input_layer = Input(shape=(MAX_SEQUENCE_LENGTH,), name="Input_Layer")
202     embedding_layer = Embedding(input_dim=weights.shape[0],
203                                output_dim=weights.shape[1],
204                                weights=[weights], trainable=False)(input_layer)
205
206     dropout_rate = hp.Choice("dropout_rate", [0.1, 0.2, 0.3])
207

```

```

208 # First BiGRU layer
209 BI_GRU_1 = tf.keras.layers.Bidirectional(GRU(units=hp.Int("GRU_unit_1",
210 min_value=64, max_value=256, step=64),
211 return_sequences=True,
212 activation="leaky_relu"))(
213 embedding_layer)
214 BI_GRU_1 = Dropout(dropout_rate)(BI_GRU_1)
215
216 # Second BiGRU layer
217 BI_GRU_2 = tf.keras.layers.Bidirectional(GRU(units=hp.Int("GRU_unit_2",
218 min_value=64, max_value=192, step=64),
219 return_sequences=True,
220 activation="leaky_relu"))(
221 BI_GRU_1)
222 BI_GRU_2 = Dropout(dropout_rate)(BI_GRU_2)
223
224 # Third BiGRU layer
225 BI_GRU_3 = tf.keras.layers.Bidirectional(GRU(units=hp.Int("GRU_unit_3",
226 min_value=64, max_value=128, step=64),
227 return_sequences=False,
228 activation="leaky_relu"))(
229 BI_GRU_2)
230 BI_GRU_3 = Dropout(dropout_rate)(BI_GRU_3)
231
232 # Dense layer
233 dense_1 = Dense(units=hp.Int("dense_unit_1", min_value=64, max_value=256,
234 step=64), activation="leaky_relu")(BI_GRU_3)
235 dense_1 = Dropout(dropout_rate)(dense_1)
236
237 # Output layer
238 output_layer = Dense(NUM_CLASSES, activation="softmax")(dense_1)
239
240 model = tf.keras.Model(inputs=input_layer, outputs=output_layer, name="
241 Text_Classification_RNN")
242
243 model.compile(optimizer="adam", loss="categorical_crossentropy", metrics=["
244 accuracy"])
245

```

```

234     return model
235
236 early_stopping = EarlyStopping(
237     monitor="val_loss",
238     patience=3,
239     restore_best_weights=True
240 )
241
242 tuner = kt.RandomSearch(
243     model_skip,
244     objective="val_loss",
245     max_trials=10,
246     executions_per_trial = 1
247 )
248
249 tuner.search(
250     X_train,
251     y_train,
252     epochs = 50,
253     batch_size = 512,
254     validation_data = (X_test, y_test),
255     callbacks=[early_stopping]
256 )
257
258 best_hyperparams_skip = tuner.get_best_hyperparameters(num_trials=1)[0]
259 best_model_skip = tuner.hypermodel.build(best_hyperparams_skip)
260 history_skip = best_model_skip.fit(X_train, y_train,
261     epochs = 50, batch_size = 512,
262     validation_data = (X_test, y_test),
263     callbacks=[early_stopping]
264 )
265
266 best_model_skip.summary()
267
268 best_hparams = tuner.get_best_hyperparameters(num_trials=1)[0]
269
270 # Print the best hyperparameters
271 for param in best_hparams.values:

```



```

272     print(f"{param}: {best_hparams.get(param)}")
273
274 from sklearn.metrics import classification_report, f1_score, precision_score,
    recall_score, confusion_matrix
275
276 # Predictions
277 y_pred_prob = best_model_skip.predict(X_test)
278 y_pred = np.argmax(y_pred_prob, axis=1)
279 y_true = np.argmax(y_test, axis=1)
280
281 print("\nClassification Report:")
282 print(classification_report(y_true, y_pred))
283
284 class_names = ["World", "Sports", "Business", "Sci/Tech"]
285 # Confusion Matrix
286 cm = confusion_matrix(y_true, y_pred)
287 plt.figure(figsize=(8,6))
288 sns.heatmap(cm, annot=True, fmt="d", cmap="Greens", xticklabels=class_names,
    yticklabels=class_names)
289
290 plt.xlabel("Predicted Class")
291 plt.ylabel("True Class")
292 plt.title("Confusion Matrix")
293 plt.show()
294
295 path_to_glove_file = "glove.6B.100d.txt"
296
297 embeddings_index = {}
298 with open(path_to_glove_file) as f:
299     for line in f:
300         word, coefs = line.split(maxsplit=1)
301         coefs = np.fromstring(coefs, "f", sep=" ")
302         embeddings_index[word] = coefs
303
304 print("Found %s word vectors." % len(embeddings_index))
305
306 embedding_dim = 100
307 vocab_size = len(vocab)

```

```

308
309 glove_embedding_matrix = np.zeros((vocab_size, embedding_dim))
310
311 for word, i in vocab.items():
312     embedding_vector = embeddings_index.get(word)
313     if embedding_vector is not None:
314         glove_embedding_matrix[i] = embedding_vector
315
316 #Model 2: Deep learning model using glove word embedding
317
318 def model_glove(hp):
319     input_layer = Input(shape=(MAX_SEQUENCE_LENGTH,), name="Input_Layer")
320     embedding_layer = Embedding(input_dim=weights.shape[0],
321                                output_dim=weights.shape[1],
322                                weights=[glove_embedding_matrix], trainable=
323                                    False)(input_layer)
324
325     dropout_rate = hp.Choice("dropout_rate", [0.1, 0.2, 0.3])
326
327     # First BiGRU layer
328     BI_GRU_1 = tf.keras.layers.Bidirectional(GRU(units=hp.Int("GRU_unit_1",
329                                                                min_value=64, max_value=256, step=64),
330                                                                return_sequences=True,
331                                                                activation="leaky_relu"))(
332         embedding_layer)
333
334     BI_GRU_1 = Dropout(dropout_rate)(BI_GRU_1)
335
336     # Second BiGRU layer
337     BI_GRU_2 = tf.keras.layers.Bidirectional(GRU(units=hp.Int("GRU_unit_2",
338                                                                min_value=64, max_value=192, step=64),
339                                                                return_sequences=True,
340                                                                activation="leaky_relu"))(
341         BI_GRU_1)
342
343     BI_GRU_2 = Dropout(dropout_rate)(BI_GRU_2)
344
345     # Third BiGRU layer
346     BI_GRU_3 = tf.keras.layers.Bidirectional(GRU(units=hp.Int("GRU_unit_3",
347                                                                min_value=64, max_value=128, step=64),

```

```

338         return_sequences=False,
339         activation="leaky_relu"))(
340         BI_GRU_2)
341
342     BI_GRU_3 = Dropout(dropout_rate)(BI_GRU_3)
343
344     # Dense layer
345     dense_1 = Dense(units=hp.Int("dense_unit_1", min_value=64, max_value=256,
346         step=64), activation="leaky_relu")(BI_GRU_3)
347     dense_1 = Dropout(dropout_rate)(dense_1)
348
349     # Output layer
350     output_layer = Dense(NUM_CLASSES, activation="softmax")(dense_1)
351
352     model = tf.keras.Model(inputs=input_layer, outputs=output_layer, name="
353         Text_Classification_RNN")
354
355     model.compile(optimizer="adam", loss="categorical_crossentropy", metrics=["
356         accuracy"])
357
358     return model
359
360     early_stopping = EarlyStopping(
361         monitor="val_loss",
362         patience=3,
363         restore_best_weights=True
364     )
365
366     tuner = kt.RandomSearch(
367         model_glove,
368         objective="val_loss",
369         max_trials=10,
370         executions_per_trial = 1
371     )
372
373     tuner.search(
374         X_train,
375         y_train,
376         epochs = 50,

```

```

371     batch_size = 512,
372     validation_data = (X_test, y_test),
373     callbacks=[early_stopping]
374 )
375
376 best_hyperparams_glove = tuner.get_best_hyperparameters(num_trials=1)[0]
377 best_model_glove = tuner.hypermodel.build(best_hyperparams_glove)
378 history_glove = best_model_glove.fit(X_train, y_train,
379                                     epochs = 50, batch_size = 512,
380                                     validation_data = (X_test, y_test),
381                                     callbacks=[early_stopping])
382
383 best_model_glove.summary()
384
385 best_hparams = tuner.get_best_hyperparameters(num_trials=1)[0]
386
387 # Print the best hyperparameters
388 for param in best_hparams.values:
389     print(f"{param}: {best_hparams.get(param)}")
390
391 # Predictions
392 y_pred_prob = best_model_glove.predict(X_test)
393 y_pred = np.argmax(y_pred_prob, axis=1)
394 y_true = np.argmax(y_test, axis=1)
395
396 print("\nClassification Report:")
397 print(classification_report(y_true, y_pred))
398
399 class_names = ["World", "Sports", "Business", "Sci/Tech"]
400 # Confusion Matrix
401 cm = confusion_matrix(y_true, y_pred)
402
403 plt.figure(figsize=(8,6))
404 sns.heatmap(cm, annot=True, fmt="d", cmap="Reds", xticklabels=class_names,
405             yticklabels=class_names)
406 plt.xlabel("Predicted Class")
407 plt.ylabel("True Class")
408 plt.title("Confusion Matrix")

```

```

408 plt.show()
409
410 # Task 2
411
412 import numpy as np
413 import pandas as pd
414 import os
415 import random
416 import re
417 import tensorflow as tf
418 import tqdm
419 from tensorflow.keras import layers
420 from collections import Counter
421 from nltk.tokenize import word_tokenize
422 from nltk.corpus import stopwords
423 from nltk.stem import WordNetLemmatizer
424 import nltk
425 from sklearn.feature_extraction.text import TfidfVectorizer
426 import scipy.cluster.hierarchy as sch
427 from sklearn.decomposition import PCA
428 import matplotlib.pyplot as plt
429 from sklearn.cluster import KMeans
430
431
432 random.seed(407)
433
434 file_path = "french-theater.txt"
435
436 with open(file_path, "r", encoding="utf-8") as f:
437     text = f.read()
438
439 #Splitting the text and sorting them into a dictionary
440 plays = re.split(r'###(\d{2})###\n', text)[1:]
441 plays_dict = {plays[i]: plays[i+1].strip().split("#####")[0] for i in range(0,
    len(plays), 2)}
442
443 #Test
444 print(f"\nPlay {1}:")

```

```

445 print(play_text[:500])
446
447 def preprocess_text(text):
448     text = text.lower()
449     text = re.sub(r"[^\w\s]", "", text) # Remove punctuation
450     tokens = word_tokenize(text)
451     tokens = [t for t in tokens if t not in stopwords.words("french")] # Remove
        stopwords
452     return tokens
453
454 document_texts = []
455 for plays in plays_dict.values():
456     document_texts.append(preprocess_text(plays))
457
458 document_texts_joined = [" ".join(tokens) for tokens in document_texts]
459
460 tfidf_vectorizer = TfidfVectorizer(max_features=5000)
461
462 tfidf_matrix = tfidf_vectorizer.fit_transform(document_texts_joined)
463 tfidf_matrix_dense = tfidf_matrix.toarray()
464
465 linkage_matrix = sch.linkage(tfidf_matrix_dense, method='ward')
466
467 plt.figure(figsize=(10, 5))
468 sch.dendrogram(linkage_matrix, labels=list(plays_dict.keys()), leaf_rotation=90)
469 plt.title("Hierarchical Clustering Dendrogram")
470 plt.xlabel("Plays")
471 plt.ylabel("Distance")
472 plt.show()
473
474 num_clusters = 2
475 kmeans = KMeans(n_clusters=num_clusters, random_state=407, n_init=10)
476
477 kmeans.fit(tfidf_matrix_dense)
478 cluster_labels = kmeans.labels_
479
480 # Clusters
481 for play, label in zip(plays_dict.keys(), cluster_labels):

```

```

482     print(f"Play {play} -> Cluster {label}")
483
484     pca = PCA(n_components=2)
485     tfidf_pca = pca.fit_transform(tfidf_matrix_dense)
486
487     plt.figure(figsize=(8,6))
488     plt.scatter(tfidf_pca[:,0], tfidf_pca[:,1], c=cluster_labels, cmap="coolwarm",
489                 edgecolors="k")
489     for i, play_num in enumerate(plays_dict.keys()):
490         plt.annotate(play_num, (tfidf_pca[i,0], tfidf_pca[i,1]))
491
492     plt.title("PCA Projection of Plays")
493     plt.xlabel("Principal Component 1")
494     plt.ylabel("Principal Component 2")
495     plt.show()

```

Appendix B: Generative AI usage

ChatGPT was used to brainstorm and to help me develop my ideas for approaching the tasks, also as an alternative to google for debugging code.